

Thanks for the previous conversation! As a next step, we'd like to invite you to a short practical exercise followed by a 45-minute walkthrough. The goal is to see how you approach a realistic first-week problem in the Innovation team, end to end.

The task

Build a working demo of an **AI agent that reviews a corporate credit memo and flags covenant risks**.

Because real credit memos are typically under NDA, **the first step is to generate your own synthetic memo**. This mirrors what v1 POC work actually looks like in the team - you rarely start with clean production data, so fabricating plausible test data is part of the very first step.

Step 1: Generate a synthetic corporate credit memo

- Roughly 10-20 pages, PDF or DOCX, close to a real one.
- A plausible mid-market or large-corporate borrower, syndicated or bilateral, your choice.
- Should include the sections a real memo would: borrower overview, transaction summary, financial analysis, covenant package (financial and non-financial), risk factors, recommendation and etc.
- Any tool. LLM-generated is fine and expected. We want to see how you prompt for it and what you check for.

Step 2: Build the agent

The agent should:

1. Ingest the memo.
2. Extract the covenant terms (financial and non-financial).
3. Identify the three highest-risk covenants and explain why, in one to two sentences each.
4. Return a structured JSON output with the covenants, risk ranking, and reasoning.

Constraints:

- Any stack: Lovable, n8n, Python, Zapier, Firebase, Supabase, Railway, direct model APIs, whatever you'd actually use.
- Any model. Pick what you'd defend.
- Working beats polished. A rough demo that runs is worth more than a beautiful deck.
- Timebox: **48 hours from receipt**. We're not looking for weekend-consuming effort. If you hit 5-6 hours across memo generation and agent build and it works, stop there.

What to send back

1. The **synthetic memo** you generated.
2. A **Loom (or similar) walkthrough, max 10 minutes**, covering:
 - o How you generated the memo and what you had to fix or check.
 - o The agent running live on your memo.
 - o The architecture: what tools, what model, how the pieces talk to each other.
 - o One trade-off you made and why.
 - o One thing you'd do differently with two more weeks.
3. The **GitHub repo, n8n export, or shared workspace link** so we can look under the hood.

What we're evaluating

We want to be transparent about the bar. This exercise is designed to answer three questions:

- Can you build (or very closely guide) a prototype yourself, at speed?
- Can you explain the integrations you chose, not just name them?
- Can you work from an ambiguous brief to something tangible fast?

We're not evaluating UI polish, prompt sophistication, presentation style, lending domain knowledge, or covenant-analysis accuracy. A confident, running, honestly-explained rough demo will score higher than a slick one you can't defend.

Follow-up

Once we have your submission, we'll book a **45-minute session** to walk through it together, drill into the choices, and discuss how you'd extend it into a real POC.

If anything about the brief is unclear, just ask. Look forward to seeing what you build. Look forward to seeing what you build.